

Assignment5

July 4, 2024

1 Assignment 5

1.1 1. Choose a REGRESSION dataset (reusing bikeshare is allowed), perform a test/train split, and build a regression model (just like in assignment 3), and calculate the

+ Training Error (MSE, MAE)

+ Testing Error (MSE, MAE)

1.2 2. Choose a CLASSIFICATION dataset (not the adult.data set, The UCI repository has many datasets as well as Kaggle), perform test/train split and create a classification model (your choice but DecisionTree is fine). Calculate

+ Accuracy

+ Confusion Matrix

+ Classification Report

1.3 3. (Bonus) See if you can improve the classification model's performance with any tricks you can think of (modify features, remove features, polynomial features)

```
[3]: import matplotlib.pyplot as plt
      %matplotlib inline
      from sklearn import linear_model, metrics
      plt.rcParams['figure.figsize'] = 20, 10
      import pandas as pd
      import numpy as np

      bikeshare_data = pd.read_csv('bikeshare_hour_count.csv')
      bikeshare_data.dropna(inplace=True)
      bikeshare_data
```

```
[3]:
```

	hour	monday	tuesday	wednesday	thursday	friday	saturday	sunday
0	0.0	21.0	34.0	43.0	47.0	51.0	89.0	106.0
1	0.1	39.0	22.0	27.0	37.0	56.0	87.0	100.0
2	0.2	31.0	24.0	26.0	42.0	50.0	98.0	77.0
3	0.3	26.0	27.0	25.0	29.0	52.0	99.0	87.0

4	0.4	19.0	24.0	29.0	29.0	50.0	98.0	69.0
..
235	23.5	36.0	65.0	60.0	94.0	80.0	93.0	28.0
236	23.6	37.0	61.0	66.0	100.0	81.0	95.0	28.0
237	23.7	30.0	42.0	49.0	80.0	101.0	105.0	27.0
238	23.8	33.0	52.0	47.0	79.0	91.0	93.0	24.0
239	23.9	34.0	33.0	48.0	65.0	105.0	111.0	23.0

[235 rows x 8 columns]

```
[4]: #Part 1
from sklearn.linear_model import LinearRegression
from sklearn import linear_model

x = bikeshare_data['hour'].values.reshape(-1,1)
y = bikeshare_data['saturday'].values.reshape(-1,1)

from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.2)
linear = linear_model.LinearRegression().fit(xtrain, ytrain)
```

```
[6]: (
    metrics.mean_squared_error(ytest, linear.predict(xtest))
)
```

[6]: 23559.85808987989

```
[8]: (
    metrics.mean_absolute_error(ytest, linear.predict(xtest))
)
```

[8]: 136.0879897318393

```
[4]: #Part 2
heart_data = pd.read_csv('heart.csv')
heart_data_golden = pd.read_csv('heart.csv')
heart_data.dropna()
heart_data.dropna()
heart_data
```

```
[4]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	52	1	0	125	212	0	1	168	0	1.0	
1	53	1	0	140	203	1	0	155	1	3.1	
2	70	1	0	145	174	0	1	125	1	2.6	
3	61	1	0	148	203	0	1	161	0	0.0	
4	62	0	0	138	294	1	1	106	0	1.9	
...	

1020	59	1	1	140	221	0	1	164	1	0.0
1021	60	1	0	125	258	0	0	141	1	2.8
1022	47	1	0	110	275	0	0	118	1	1.0
1023	50	0	0	110	254	0	0	159	0	0.0
1024	54	1	0	120	188	0	1	113	0	1.4

	slope	ca	thal	target
0	2	2	3	0
1	0	0	3	0
2	0	0	3	0
3	2	1	3	0
4	1	3	2	0
...
1020	2	0	2	1
1021	1	1	3	0
1022	1	1	2	0
1023	2	0	2	1
1024	1	1	3	0

[1025 rows x 14 columns]

```
[5]: heart_data['oldpeak'].unique()
```

```
[5]: array([1. , 3.1, 2.6, 0. , 1.9, 4.4, 0.8, 3.2, 1.6, 3. , 0.7, 4.2, 1.5,
          2.2, 1.1, 0.3, 0.4, 0.6, 3.4, 2.8, 1.2, 2.9, 3.6, 1.4, 0.2, 2. ,
          5.6, 0.9, 1.8, 6.2, 4. , 2.5, 0.5, 0.1, 2.1, 2.4, 3.8, 2.3, 1.3,
          3.5])
```

```
[6]: non_numeric_columns = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'thal']
```

```
[7]: x = heart_data.copy().drop(non_numeric_columns, axis=1)
```

```
[8]: xt = heart_data_golden.copy().drop(non_numeric_columns, axis=1)
```

```
[9]: x['oldpeak'] = x.oldpeak.str.contains('.').astype(int)
```

```
-----
AttributeError                                Traceback (most recent call last)
Cell In[9], line 1
----> 1 x['oldpeak'] = x.oldpeak.str.contains('.').astype(int)

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
pandas/core/generic.py:5989, in NDFrame.__getattr__(self, name)
    5982 if (
    5983     name not in self._internal_names_set
    5984     and name not in self._metadata
    5985     and name not in self._accessors
```

```

5986         and self._info_axis._can_hold_identifiers_and_holds_name(name)
5987     ):
5988         return self[name]
-> 5989 return object.__getattr__(self, name)

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/accessor.py:224, in CachedAccessor.__get__(self, obj, cls)
    221 if obj is None:
    222     # we're accessing the attribute of the class, i.e., Dataset.geo
    223     return self._accessor
--> 224 accessor_obj = self._accessor(obj)
    225 # Replace the property with the accessor object. Inspired by:
    226 # https://www.pydanny.com/cached-property.html
    227 # We need to use object.__setattr__ because we overwrite __setattr__ on
    228 # NDFrame
    229 object.__setattr__(obj, self._name, accessor_obj)

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/strings/accessor.py:181, in StringMethods.__init__(self, data)
    178 def __init__(self, data) -> None:
    179     from pandas.core.arrays.string_ import StringDtype
--> 181     self._inferred_dtype = self._validate(data)
    182     self._is_categorical = is_categorical_dtype(data.dtype)
    183     self._is_string = isinstance(data.dtype, StringDtype)

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/strings/accessor.py:235, in StringMethods._validate(data)
    232 inferred_dtype = lib.infer_dtype(values, skipna=True)
    234 if inferred_dtype not in allowed_types:
--> 235     raise AttributeError("Can only use .str accessor with string values
↳ ")
    236 return inferred_dtype

AttributeError: Can only use .str accessor with string values!

```

```
[51]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
```

```
[53]: model = DecisionTreeClassifier(criterion='entropy')
```

```
[55]: model.fit(x.drop(['oldpeak'], axis=1), x.oldpeak)
```

```

-----
KeyError                                Traceback (most recent call last)
Cell In[55], line 1
----> 1 model.fit(x.drop(['species'], axis=1), x.species)

```

```

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/frame.py:5258, in DataFrame.drop(self, labels, axis, index,
↳ columns, level, inplace, errors)
    5110 def drop(
    5111     self,
    5112     labels: IndexLabel = None,
    (...)
    5119     errors: IgnoreRaise = "raise",
    5120 ) -> DataFrame | None:
    5121     """
    5122     Drop specified labels from rows or columns.
    5123     (...)
    5256         weight 1.0      0.8
    5257     """
-> 5258     return super().drop(
    5259         labels=labels,
    5260         axis=axis,
    5261         index=index,
    5262         columns=columns,
    5263         level=level,
    5264         inplace=inplace,
    5265         errors=errors,
    5266     )

```

```

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/generic.py:4549, in NDFrame.drop(self, labels, axis, index,
↳ columns, level, inplace, errors)
    4547 for axis, labels in axes.items():
    4548     if labels is not None:
-> 4549         obj = obj._drop_axis(labels, axis, level=level, errors=errors)
    4551 if inplace:
    4552     self._update_inplace(obj)

```

```

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/generic.py:4591, in NDFrame._drop_axis(self, labels, axis, level,
↳ errors, only_slice)
    4589     new_axis = axis.drop(labels, level=level, errors=errors)
    4590     else:
-> 4591     new_axis = axis.drop(labels, errors=errors)
    4592     indexer = axis.get_indexer(new_axis)
    4594 # Case for non-unique axis
    4595 else:

```

```

File /opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-packages/
↳ pandas/core/indexes/base.py:6699, in Index.drop(self, labels, errors)
    6697 if mask.any():
    6698     if errors != "ignore":

```

```
-> 6699         raise KeyError(f"{list(labels[mask])} not found in axis")
    6700     indexer = indexer[~mask]
    6701     return self.delete(indexer)
```

```
KeyError: "['species'] not found in axis"
```

```
[ ]: list(zip(x.drop(['oldpeak'], axis=1).columns, model.feature_importances_))
```

```
[ ]: x.drop(['oldpeak'], axis=1).head()
```

```
[ ]: set(x.columns) - set(xt.columns)
```

```
[ ]: predictions = model.predict(xt.drop(['oldpeak'], axis=1))
```

```
[ ]: predictions_train = model.predict(x.drop(['oldpeak'], axis=1))
```

```
[ ]: xt.drop(['oldpeak'], axis=1).head()
```

```
[ ]: # References
    Janosi,Andras, Steinbrunn, William, Pfisterer, Matthias,
    and Detrano,Robert.
    (1988).
    Heart Disease.
    UCI Machine Learning Repository.
    https://doi.org/10.24432/C52P4X.
```