

Assignment6

July 10, 2024

1 Assignment is below at the end

- <https://scikit-learn.org/stable/modules/tree.html>
- <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>
- https://scikit-learn.org/stable/modules/generated/sklearn.tree.plot_tree.html

```
[1]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 6)
plt.rcParams['font.size'] = 14
import pandas as pd
```

```
[2]: df = pd.read_csv('adult.data', index_col=False)
```

```
[3]: golden = pd.read_csv('adult.test', index_col=False)
```

```
[4]: golden.head()
```

```
[4]:   age  workclass  fnlwgt   education  education-num   marital-status \
0   25   Private  226802      11th              7   Never-married
1   38   Private   89814   HS-grad              9  Married-civ-spouse
2   28  Local-gov  336951  Assoc-acdm             12  Married-civ-spouse
3   44   Private  160323  Some-college            10  Married-civ-spouse
4   18      ?   103497  Some-college            10   Never-married
```

```
   occupation  relationship   race   sex  capital-gain \
0  Machine-op-inspct   Own-child  Black  Male           0
1   Farming-fishing   Husband  White  Male           0
2   Protective-serv   Husband  White  Male           0
3  Machine-op-inspct   Husband  Black  Male       7688
4      ?   Own-child  White  Female           0
```

```
   capital-loss  hours-per-week  native-country  salary
0           0           40  United-States  <=50K.
1           0           50  United-States  <=50K.
2           0           40  United-States  >50K.
3           0           40  United-States  >50K.
```

```
4          0          30  United-States  <=50K.
```

```
[5]: df.head()
```

```
[5]:   age      workclass  fnlwgt  education  education-num \
0    39      State-gov   77516   Bachelors             13
1    50  Self-emp-not-inc   83311   Bachelors             13
2    38        Private  215646    HS-grad              9
3    53        Private  234721      11th              7
4    28        Private  338409   Bachelors             13

      marital-status      occupation  relationship    race    sex \
0      Never-married      Adm-clerical  Not-in-family  White  Male
1  Married-civ-spouse  Exec-managerial      Husband  White  Male
2        Divorced  Handlers-cleaners  Not-in-family  White  Male
3  Married-civ-spouse  Handlers-cleaners      Husband  Black  Male
4  Married-civ-spouse  Prof-specialty      Wife    Black  Female

      capital-gain  capital-loss  hours-per-week  native-country  salary
0          2174           0           40  United-States  <=50K
1           0           0           13  United-States  <=50K
2           0           0           40  United-States  <=50K
3           0           0           40  United-States  <=50K
4           0           0           40          Cuba  <=50K
```

```
[6]: df.columns
```

```
[6]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
      'marital-status', 'occupation', 'relationship', 'race', 'sex',
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
      'salary'],
      dtype='object')
```

```
[7]: from sklearn import preprocessing
```

```
[8]: # Columns we want to transform
transform_columns = ['sex']

#Columns we can't use because non-numerical
non_num_columns = ['workclass', 'education', 'marital-status',
      'occupation', 'relationship', 'race', 'sex',
      'native-country']
```

1.1 First let's try using `pandas.get_dummies()` to transform columns

```
[9]: dummies = pd.get_dummies(df[transform_columns])  
dummies
```

```
[9]:      sex_ Female  sex_ Male  
0          False      True  
1          False      True  
2          False      True  
3          False      True  
4           True     False  
...         ...      ...  
32556        True     False  
32557        False      True  
32558        True     False  
32559        False      True  
32560        True     False
```

```
[32561 rows x 2 columns]
```

```
[10]: dummies.shape
```

```
[10]: (32561, 2)
```

1.2 sklearn has a similar process for OneHot Encoding features

```
[11]: onehot = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist",  
      ↪sparse=False)  
onehot.fit(df[transform_columns])
```

```
/opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-  
packages/sklearn/preprocessing/_encoders.py:972: FutureWarning: `sparse` was  
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.  
`sparse_output` is ignored unless you leave `sparse` to its default value.  
warnings.warn(
```

```
[11]: OneHotEncoder(handle_unknown='infrequent_if_exist', sparse=False,  
      sparse_output=False)
```

```
[12]: onehot.categories_
```

```
[12]: [array([' Female', ' Male'], dtype=object)]
```

```
[13]: sex = onehot.transform(df[transform_columns])  
sex
```

```
[13]: array([[0., 1.],
           [0., 1.],
           [0., 1.],
           ...,
           [1., 0.],
           [0., 1.],
           [1., 0.]])
```

```
[14]: sex.shape
```

```
[14]: (32561, 2)
```

1.3 In addition to OneHot encoding there is Ordinal Encoding

```
[15]: enc = preprocessing.OrdinalEncoder()
      enc.fit(df[["salary"]])
      salary = enc.transform(df[["salary"]])
      salary
```

```
[15]: array([[0.],
           [0.],
           [0.],
           ...,
           [0.],
           [0.],
           [1.]])
```

```
[16]: enc.categories_[0]
```

```
[16]: array([' <=50K', ' >50K'], dtype=object)
```

```
[61]: x = df.copy()

      # transformed = pd.get_dummies(df[transform_columns])

      #onehot = preprocessing.OneHotEncoder(handle_unknown="infrequent_if_exist",
      ↪sparse=False).fit(df[transform_columns])

      enc = preprocessing.OrdinalEncoder()
      enc.fit(df[["salary"]])

      x = df.copy()
      x = pd.concat([x.drop(non_num_columns, axis=1),
                    pd.get_dummies(df[transform_columns]), axis=1)
```

```

#transformed = onehot.transform(df[transform_columns])
#new_cols = list(onehot.categories_[0].flatten())
#df_trans = pd.DataFrame(transformed, columns=new_cols)

#x = pd.concat(
#    [
#        x.drop(non_num_columns, axis=1),
#        df_trans
#    ],
#    axis=1,)

x["salary"] = enc.transform(df[["salary"]])

```

[62]: x.head()

```

[62]:   age  fnlwgt  education-num  capital-gain  capital-loss  hours-per-week  \
0    39   77516             13           2174             0             40
1    50   83311             13              0             0             13
2    38  215646              9              0             0             40
3    53  234721              7              0             0             40
4    28  338409             13              0             0             40

      salary  sex_ Female  sex_ Male
0         0.0         False       True
1         0.0         False       True
2         0.0         False       True
3         0.0         False       True
4         0.0          True      False

```

[63]: xt = golden.copy()

```

transformed = onehot.transform(xt[transform_columns])
new_cols = list(onehot.categories_[0].flatten())
#df_trans = pd.DataFrame(transformed, columns=new_cols)

#x = pd.concat(
#    [
#        xt.drop(non_num_columns, axis=1),
#        df_trans
#    ],
#    axis=1,)
# xt["salary"] = enc.fit_transform(golden[["salary"]])
xt = pd.concat([xt.drop(non_num_columns, axis=1),
                pd.get_dummies(xt[transform_columns])], axis=1)
xt["salary"] = enc.fit_transform(xt[["salary"]])

```

```
[64]: xt.salary.value_counts()
```

```
[64]: salary
      0.0    12435
      1.0     3846
      Name: count, dtype: int64
```

```
[65]: enc.categories_
```

```
[65]: [array(['<=50K.', '>50K.'], dtype=object)]
```

```
[66]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.ensemble import GradientBoostingClassifier
```

Choose the model of your preference: DecisionTree or RandomForest

```
[67]: model = RandomForestClassifier(criterion='entropy')
```

```
[68]: model = DecisionTreeClassifier(criterion='entropy', max_depth=None)
```

```
[69]: model.fit(x.drop(['fnlwgt', 'salary'], axis=1), x.salary)
```

```
[69]: DecisionTreeClassifier(criterion='entropy')
```

```
[70]: model.tree_.node_count
```

```
[70]: 8317
```

```
[71]: list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.
      ↪feature_importances_))
```

```
[71]: [('age', 0.32366222896975533),
      ('education-num', 0.15861240773791468),
      ('capital-gain', 0.2284592742844192),
      ('capital-loss', 0.07843149812577388),
      ('hours-per-week', 0.15529640153538432),
      ('sex_ Female', 0.021232345235427594),
      ('sex_ Male', 0.03430584411132491)]
```

```
[72]: list(zip(x.drop(['fnlwgt', 'salary'], axis=1).columns, model.
      ↪feature_importances_))
```

```
[72]: [('age', 0.32366222896975533),
      ('education-num', 0.15861240773791468),
      ('capital-gain', 0.2284592742844192),
      ('capital-loss', 0.07843149812577388),
```

```
( 'hours-per-week', 0.15529640153538432),
( 'sex_ Female', 0.021232345235427594),
( 'sex_ Male', 0.03430584411132491)]
```

```
[73]: x.drop(['fnlwgt', 'salary'], axis=1).head()
```

```
[73]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week	\
0	39	13	2174	0	40	
1	50	13	0	0	13	
2	38	9	0	0	40	
3	53	7	0	0	40	
4	28	13	0	0	40	

	sex_ Female	sex_ Male
0	False	True
1	False	True
2	False	True
3	False	True
4	True	False

```
[74]: set(x.columns) - set(xt.columns)
```

```
[74]: set()
```

```
[75]: list(x.drop('salary', axis=1).columns)
```

```
[75]: ['age',
      'fnlwgt',
      'education-num',
      'capital-gain',
      'capital-loss',
      'hours-per-week',
      'sex_ Female',
      'sex_ Male']
```

```
[76]: predictions = model.predict(xt.drop(['fnlwgt', 'salary'], axis=1))
      predictionsx = model.predict(x.drop(['fnlwgt', 'salary'], axis=1))
```

```
[77]: from sklearn.metrics import (
      accuracy_score,
      classification_report,
      confusion_matrix, auc, roc_curve
      )
```

```
[78]: accuracy_score(xt.salary, predictions)
```

```
[78]: 0.821018364965297
```

```
[79]: accuracy_score(xt.salary, predictions)
```

```
[79]: 0.821018364965297
```

```
[80]: confusion_matrix(xt.salary, predictions)
```

```
[80]: array([[11462,  973],
        [ 1941, 1905]])
```

```
[81]: print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.86	0.92	0.89	12435
1.0	0.66	0.50	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
[82]: print(classification_report(xt.salary, predictions))
```

	precision	recall	f1-score	support
0.0	0.86	0.92	0.89	12435
1.0	0.66	0.50	0.57	3846
accuracy			0.82	16281
macro avg	0.76	0.71	0.73	16281
weighted avg	0.81	0.82	0.81	16281

```
[83]: accuracy_score(x.salary, predictionsx)
```

```
[83]: 0.8955806025613464
```

```
[84]: confusion_matrix(x.salary, predictionsx)
```

```
[84]: array([[24097,  623],
        [ 2777, 5064]])
```

```
[85]: print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841

accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

```
[86]: print(classification_report(x.salary, predictionsx))
```

	precision	recall	f1-score	support
0.0	0.90	0.97	0.93	24720
1.0	0.89	0.65	0.75	7841

accuracy			0.90	32561
macro avg	0.89	0.81	0.84	32561
weighted avg	0.90	0.90	0.89	32561

2 For the following use the above adult dataset.

3 1. Show the RandomForest outperforms the DecisionTree for a fixed max_depth by training using the train set and calculate precision, recall, f1, confusion matrix on golden-test set. Start with only numerical features/columns. (age, education-num, capital-gain, capital-loss, hours-per-week)

```
[95]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import precision_score, recall_score, f1_score,
      ↪confusion_matrix
import pandas as pd

# Assuming train_set and golden_test_set are DataFrames containing the required
      ↪columns

# Selecting only the numerical features
numerical_features = ['age', 'education-num', 'capital-gain', 'capital-loss',
      ↪'hours-per-week']

# Train set
X_train = x[numerical_features]
y_train = x['salary'] # Replace 'target_column' with the actual target column
      ↪name

# Golden-test set
```

```

X_test = xt[numerical_features]
y_test = xt['salary'] # Replace 'target_column' with the actual target column
                        ↳ name

# Train a DecisionTree model
dt_model = DecisionTreeClassifier(max_depth=5) # Using a fixed max_depth of 5
dt_model.fit(X_train, y_train)

# Train a RandomForest model
rf_model = RandomForestClassifier(max_depth=5, n_estimators=100) # Using a
                        ↳ fixed max_depth of 5 and 100 trees
rf_model.fit(X_train, y_train)

# Making predictions
dt_predictions = dt_model.predict(X_test)
rf_predictions = rf_model.predict(X_test)

# Calculate precision, recall, and F1 score
dt_precision = precision_score(y_test, dt_predictions)
dt_recall = recall_score(y_test, dt_predictions)
dt_f1 = f1_score(y_test, dt_predictions)

rf_precision = precision_score(y_test, rf_predictions)
rf_recall = recall_score(y_test, rf_predictions)
rf_f1 = f1_score(y_test, rf_predictions)

# Calculate confusion matrix
dt_conf_matrix = confusion_matrix(y_test, dt_predictions)
rf_conf_matrix = confusion_matrix(y_test, rf_predictions)

```

- 4 2. Use a RandomForest or DecisionTree and the adult dataset, systematically add new columns, one by one, that are non-numerical but converted using the feature-extraction techniques we learned. Using the golden-test set show [precision, recall, f1, confusion matrix] for each additional feature added.

```

[103]: import pandas as pd
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.metrics import precision_score, recall_score, f1_score,
        ↳ confusion_matrix

```

```

[113]: non_numerical_columns = ['workclass', 'education', 'marital-status',
        ↳ 'occupation', 'relationship', 'race', 'sex', 'native-country']

```

```

xt = pd.get_dummies(xt, columns=['salary'])

for column in non_numerical_columns:
    rf_model = RandomForestClassifier(max_depth=5, n_estimators=100)
    rf_model.fit(X_train, y_train)
    rf_predictions = rf_model.predict(X_test)
    rf_precision = precision_score(y_test, rf_predictions)
    rf_recall = recall_score(y_test, rf_predictions)
    rf_f1 = f1_score(y_test, rf_predictions)
    rf_conf_matrix = confusion_matrix(y_test, rf_predictions)
    print(f"Performance metrics for {column}: Precision={rf_precision},  

    Recall={rf_recall}, F1={rf_f1}, Confusion Matrix={rf_conf_matrix}")

```

```

Performance metrics for workclass: Precision=0.7644529383659818,
Recall=0.4160166406656266, F1=0.5388112476847955, Confusion Matrix=[[11942
493]
 [ 2246  1600]]
Performance metrics for education: Precision=0.763445978105664,
Recall=0.4170566822672907, F1=0.5394316462081722, Confusion Matrix=[[11938
497]
 [ 2242  1604]]
Performance metrics for marital-status: Precision=0.7629453681710214,
Recall=0.4175767030681227, F1=0.5397412199630315, Confusion Matrix=[[11936
499]
 [ 2240  1606]]
Performance metrics for occupation: Precision=0.7628326996197718,
Recall=0.4173166926677067, F1=0.5394957983193277, Confusion Matrix=[[11936
499]
 [ 2241  1605]]
Performance metrics for relationship: Precision=0.7735294117647059,
Recall=0.41029641185647425, F1=0.5361875637104995, Confusion Matrix=[[11973
462]
 [ 2268  1578]]
Performance metrics for race: Precision=0.7653256704980843,
Recall=0.4154966198647946, F1=0.5385911695315133, Confusion Matrix=[[11945
490]
 [ 2248  1598]]
Performance metrics for sex: Precision=0.7722195240407965,
Recall=0.41341653666146644, F1=0.5385266723116003, Confusion Matrix=[[11966
469]
 [ 2256  1590]]
Performance metrics for native-country: Precision=0.7601142313184198,
Recall=0.41523660946437857, F1=0.5370775180763411, Confusion Matrix=[[11931
504]
 [ 2249  1597]]

```