

Assignment7

July 10, 2024

1 Assignment is at the bottom!

```
[1]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np


from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model
```

```
[2]: y = np.concatenate([np.zeros(10), np.ones(10)])
x = np.linspace(0, 10, len(y))
```

```
[3]: plt.scatter(x, y, c=y)
```

```
[3]: <matplotlib.collections.PathCollection at 0x7fe84e507b50>
```




```
[4]: model = LogisticRegression()
```

```
[5]: model.fit(x.reshape(-1, 1), y)
```

```
[5]: LogisticRegression()
```

```
[6]: plt.scatter(x, y, c=y)  
plt.plot(x, model.predict_proba(x.reshape(-1, 1))[:, 1])
```

```
[6]: [<matplotlib.lines.Line2D at 0x7fe844e1a250>]
```




```
[7]: b, b0 = model.coef_, model.intercept_  
model.coef_, model.intercept_
```

```
[7]: (array([[1.46709085]]), array([-7.33542562]))
```

```
[8]: plt.plot(x, 1/(1+np.exp(-x)))
```

```
[8]: [<matplotlib.lines.Line2D at 0x7fe844ec5410>]
```




```
[9]: b
```

```
[9]: array([[1.46709085]])
```

```
[10]: plt.plot(x, 1/(1+np.exp(-(b[0]*x +b0))))
```


```
[10]: [<matplotlib.lines.Line2D at 0x7fe84dfc9090>]
```



```
[19]: from mpl_toolkits import mplot3d
      from mpl_toolkits.mplot3d import Axes3D
      import matplotlib.pyplot as plt
      from matplotlib import cm
      from matplotlib.ticker import LinearLocator, FormatStrFormatter
      import numpy as np

      fig = plt.figure()
      ax = fig.add_subplot(projection='3d')

      # Make data.
      X = np.arange(-10, 10, 0.25)
      Y = np.arange(-10, 10, 0.25)
      X, Y = np.meshgrid(X, Y)
      R = np.sqrt(X**2 + Y**2)
      Z = 1/(1+np.exp(-(b[0]*X + b[0]*Y + b0)))
      surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                             linewidth=0, antialiased=False)
```



[20]: X

```
[20]: array([[ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
            [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
            [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
            ...,
            [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
            [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
            [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75]])
```

[21]: Y


```
[21]: array([[ -10.   ,  -10.   ,  -10.   , ...,  -10.   ,  -10.   ,  -10.   ],
            [  -9.75,  -9.75,  -9.75, ...,  -9.75,  -9.75,  -9.75],
            [  -9.5 ,  -9.5 ,  -9.5 , ...,  -9.5 ,  -9.5 ,  -9.5 ],
            ...,
            [   9.25,   9.25,   9.25, ...,   9.25,   9.25,   9.25],
            [   9.5 ,   9.5 ,   9.5 , ...,   9.5 ,   9.5 ,   9.5 ],
            [   9.75,   9.75,   9.75, ...,   9.75,   9.75,   9.75]])
```

What if the data doesn't really fit this pattern?

```
[22]: y = np.concatenate([np.zeros(10), np.ones(10), np.zeros(10)])
      x = np.linspace(0, 10, len(y))
```

```
[23]: plt.scatter(x,y, c=y)
```

```
[23]: <matplotlib.collections.PathCollection at 0x7fe8449fc650>
```




```
[24]: model.fit(x.reshape(-1, 1),y)
```

```
[24]: LogisticRegression()
```

```
[25]: plt.scatter(x,y)
      plt.plot(x, model.predict_proba(x.reshape(-1, 1)))
```

```
[25]: [<matplotlib.lines.Line2D at 0x7fe844b113d0>,
      <matplotlib.lines.Line2D at 0x7fe842f1d1d0>]
```



```
[26]: model1 = LogisticRegression()
      model1.fit(x[:15].reshape(-1, 1),y[:15])
```


```
[26]: LogisticRegression()
```

```
[27]: model2 = LogisticRegression()
      model2.fit(x[15:].reshape(-1, 1),y[15:])
```

```
[27]: LogisticRegression()
```

```
[28]: plt.scatter(x,y, c=y)
      plt.plot(x, model1.predict_proba(x.reshape(-1, 1))[:,1] * model2.
      ↪predict_proba(x.reshape(-1, 1))[:,1])
```

```
[28]: [<matplotlib.lines.Line2D at 0x7fe842f1e650>]
```



```
[30]: df = pd.read_csv('adult.data', index_col=False)
golden = pd.read_csv('adult.test', index_col=False)
```

```
[31]: from sklearn import preprocessing

enc = preprocessing.OrdinalEncoder()
```

```
[32]: transform_columns = ['sex', 'workclass', 'education', 'marital-status',
                           'occupation', 'relationship', 'race', 'sex',
                           'native-country', 'salary']
```

```
[33]: x = df.copy()

x[transform_columns] = enc.fit_transform(df[transform_columns])

golden['salary'] = golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', '
↳ >50K')
xt = golden.copy()

xt[transform_columns] = enc.transform(golden[transform_columns])
```

```
[34]: df.salary.unique()
```

```
[34]: array([' <=50K', ' >50K'], dtype=object)
```

```
[35]: golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K').unique()
```

```
[35]: array([' <=50K', ' >50K'], dtype=object)
```



```
[36]: model.fit(preprocessing.scale(x.drop('salary', axis=1)), x.salary)
```

```
[36]: LogisticRegression()
```

```
[37]: pred = model.predict(preprocessing.scale(x.drop('salary', axis=1)))  
pred_test = model.predict(preprocessing.scale(xt.drop('salary', axis=1)))
```

```
[38]: x.head()
```

```
[38]:
```

	age	workclass	fnlwtg	education	education-num	marital-status	\
0	39	7.0	77516	9.0	13	4.0	
1	50	6.0	83311	9.0	13	2.0	
2	38	4.0	215646	11.0	9	0.0	
3	53	4.0	234721	1.0	7	2.0	
4	28	4.0	338409	9.0	13	2.0	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	1.0	1.0	4.0	1.0	2174	0	
1	4.0	0.0	4.0	1.0	0	0	
2	6.0	1.0	4.0	1.0	0	0	
3	6.0	0.0	2.0	1.0	0	0	
4	10.0	5.0	2.0	0.0	0	0	

	hours-per-week	native-country	salary
0	40	39.0	0.0
1	13	39.0	0.0
2	40	39.0	0.0
3	40	39.0	0.0
4	40	5.0	0.0

```
[39]: from sklearn.metrics import (  
      accuracy_score,  
      classification_report,  
      confusion_matrix, auc, roc_curve  
      )
```

```
[40]: accuracy_score(x.salary, pred)
```

```
[40]: 0.8250360861152913
```

```
[41]: confusion_matrix(x.salary, pred)
```

```
[41]: array([[23300, 1420],  
       [ 4277, 3564]])
```

```
[42]: print(classification_report(x.salary, pred))
```

	precision	recall	f1-score	support
0.0	0.84	0.94	0.89	24720
1.0	0.72	0.45	0.56	7841
accuracy			0.83	32561
macro avg	0.78	0.70	0.72	32561
weighted avg	0.81	0.83	0.81	32561

```
[43]: print(classification_report(xt.salary, pred_test))
```

	precision	recall	f1-score	support
0.0	0.85	0.94	0.89	12435
1.0	0.70	0.45	0.55	3846
accuracy			0.82	16281
macro avg	0.77	0.69	0.72	16281
weighted avg	0.81	0.82	0.81	16281

2 Assignment

- 2.1 1. Use your own dataset (Heart.csv is acceptable), create a train and a test set, and build 2 models: Logistic Regression and Decision Tree (shallow). Compare the test results using `classification_report` and `confusion_matrix`. Explain which algorithm is optimal
- 2.2 2. Repeat 1. but let the Decision Tree be much deeper to allow over-fitting. Compare the two models' test results again, and explain which is optimal

```
[47]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler

# Load the dataset
heart_data = pd.read_csv('heart.csv')

# Split the dataset into features and target variable
X = heart_data.drop('target', axis=1)
y = heart_data['target']

# Split the dataset into a training set and a test set
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Build a Logistic Regression model
logistic_model = LogisticRegression()
logistic_model.fit(X_train, y_train)

# Build a Decision Tree model
tree_model = DecisionTreeClassifier(max_depth=3)
tree_model.fit(X_train, y_train)

# Evaluate the Logistic Regression model
logistic_predictions = logistic_model.predict(X_test)
print("Logistic Regression Model:")
print(classification_report(y_test, logistic_predictions))
print(confusion_matrix(y_test, logistic_predictions))

# Evaluate the Decision Tree model
tree_predictions = tree_model.predict(X_test)
print("Decision Tree Model:")
print(classification_report(y_test, tree_predictions))
print(confusion_matrix(y_test, tree_predictions))

```

Logistic Regression Model:

	precision	recall	f1-score	support
0	0.86	0.70	0.77	102
1	0.75	0.88	0.81	103
accuracy			0.79	205
macro avg	0.80	0.79	0.79	205
weighted avg	0.80	0.79	0.79	205

```

[[71 31]
 [12 91]]

```

Decision Tree Model:

	precision	recall	f1-score	support
0	0.85	0.68	0.75	102
1	0.73	0.88	0.80	103
accuracy			0.78	205
macro avg	0.79	0.78	0.78	205
weighted avg	0.79	0.78	0.78	205

```

[[69 33]
 [12 91]]

```

```
/opt/conda/envs/anaconda-panel-2023.05-py310/lib/python3.11/site-  
packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
[50]: # In this case, it appears that the Logistic Regression  
# Model is the slightly superior model, because its f1-score  
# and accuracy are higher among other metrics.
```

```
[49]: from sklearn.metrics import accuracy_score, precision_score, recall_score,   
      ↪ f1_score  
  
      # Split the dataset into features and target variable  
      X = heart_data.drop('target', axis=1)  
      y = heart_data['target']  
  
      # Split the data into training and testing sets  
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,   
      ↪ random_state=42)  
  
      # Train the original decision tree model  
      original_model = DecisionTreeClassifier(max_depth=3) # Original depth  
      original_model.fit(X_train, y_train)  
  
      # Test the original decision tree model  
      original_y_pred = original_model.predict(X_test)  
  
      # Evaluate the performance of the original model  
      original_accuracy = accuracy_score(y_test, original_y_pred)  
      original_precision = precision_score(y_test, original_y_pred)  
      original_recall = recall_score(y_test, original_y_pred)  
      original_f1 = f1_score(y_test, original_y_pred)  
  
      # Train the overfit decision tree model  
      overfit_model = DecisionTreeClassifier(max_depth=10) # Much deeper depth for  
      ↪ overfitting  
      overfit_model.fit(X_train, y_train)  
  
      # Test the overfit decision tree model  
      overfit_y_pred = overfit_model.predict(X_test)
```

```

# Evaluate the performance of the overfit model
overfit_accuracy = accuracy_score(y_test, overfit_y_pred)
overfit_precision = precision_score(y_test, overfit_y_pred)
overfit_recall = recall_score(y_test, overfit_y_pred)
overfit_f1 = f1_score(y_test, overfit_y_pred)

# Compare the test results of the two models
print("Original Decision Tree Model:")
print("Accuracy:", original_accuracy)
print("Precision:", original_precision)
print("Recall:", original_recall)
print("F1 Score:", original_f1)

print("\nOverfit Decision Tree Model:")
print("Accuracy:", overfit_accuracy)
print("Precision:", overfit_precision)
print("Recall:", overfit_recall)
print("F1 Score:", overfit_f1)

```

Original Decision Tree Model:
Accuracy: 0.7804878048780488
Precision: 0.7338709677419355
Recall: 0.883495145631068
F1 Score: 0.8017621145374451

Overfit Decision Tree Model:
Accuracy: 0.9853658536585366
Precision: 1.0
Recall: 0.970873786407767
F1 Score: 0.9852216748768473

```

[51]: # In this case, the 'Overfit' Decision Tree Model
# outclasses the Logistic Regression Model as its accuracy
# is almost perfect and its precision actually is 1.0!
# Furthermore, the F1 Score is extremely high.

```